

AI Presentation Coach

Using AI to make you a better public speaker

An AWS-native event based agentic presentation coaching platform built with Kiro and AI-DLC practices.

This is not just a demo, it is a useful pattern for taking agentic systems from idea to PoC to Production via governed software delivery.

Created and developed by Michael Geiser

Dark theme retained from the working deck



AI Presentation Coach

Using AI to make you a better public speaker

AI-Powered Presentation Coaching

Upload your presentation audio and receive detailed feedback across 7 dimensions: delivery, structure, executive presence, technical communication, audience engagement, pacing, and persuasion.

[Sign In to Get Started](#)

Presentation Coaching Report

User: Michael Geiser (mgeiser@mgeiser.net)
Presentation Title: new test
Description: test of Exec summary
Presentation File: 1_Introducing AgentCore.mp3
Uploaded: June 22, 2025 21:42 ET

Executive Summary

This presentation is technically knowledgeable and genuinely engaging in tone, but it reads more like an enthusiastic product walkthrough than a confident, authoritative introduction — and that gap is what the 5.9 overall average reflects. That score signals a presenter with real substance who has not yet learned to carry that substance with full conviction. The content is credible, the energy is real, but the delivery patterns, structural looseness, and absence of persuasive architecture consistently blunt the impact of what is actually strong material.

The clearest strengths here are the technical depth and the instinct to make complex infrastructure concepts feel human. Delivery scores the highest of any dimension for good reason: the conversational warmth, inclusive second-person framing, and natural rhetorical questions in the memory section all create a sense of a knowledgeable guide rather than a lecturer reading slides. Technical Communication also scores well because the problem-first framing is genuinely effective — introducing each capability by naming the pain point it solves is sound instructional architecture. The honest acknowledgment of platform limitations and the proactive handling of the vendor-lock objection demonstrate a presenter who understands the audience's skepticism, and that kind of intellectual transparency builds real credibility with technical listeners.

The two highest-leverage improvements cut across nearly every lower-scoring dimension simultaneously. First, Executive Presence and Persuasion both suffer from the same root cause: hedging and casual filler language — “kind of,” “sort of,” “super super easy,” “mash,” “if you ask me” — that systematically undercuts the authority this presenter has clearly earned. Replacing that language with precise, confident phrasing would lift scores across Delivery, Executive Presence, and Persuasion in a single pass. Second, Audience Engagement and Structure both suffer from the absence of a narrative spine. The presentation is currently a feature inventory with a loose framing device, transforming it into a problem-solution arc with a concrete protagonist — even a hypothetical developer struggling to get an agent into production — would give the audience something to follow emotionally, not just intellectually. Structure also needs a clean verbal agenda at the opening and a proper 30-to-45-second closing recap, rather than the current 12-second pivot to the demo. For the next recording, focus on three things. Open with one bold, evidence-anchored statement of impact instead of “hey everyone,” replace every instance of “kind of” and “sort of” with either a precise word or a clean pause, and close with a 40-second summary that names the three capabilities the audience should remember and tells them exactly what they are about to do in the hands-on session and why it matters.

**AWS-native
serverless +
agentic**

**Built for
zero idle cost**

EXECUTIVE LENS

The CTO readout

The AI Presentation Coach ("PresCoach") is a working AWS-native product pattern, not a lab experiment. It accepts audio, prepares it, then hands it to independent Agentic AI evaluators that produce a coaching report.

7

AI coaching dimensions

20+AWS native
services used**726**

Tests through all workflows

\$33.69Spend for environment
and AI tools used

(\$6.59 AWS + \$27.10 Kiro)

**Architecture decision
worth a CTO review****Where to harden
before real users****How to scale
without cost traps****What AI-DLC
actually changed**

WHAT IT DOES

Using the product is simple, the system behind it is not

The screenshot shows the 'Upload Presentation' interface. At the top, there's a navigation bar with 'Presentation Coach' and 'Administrator' (highlighted in red), along with links for 'Upload', 'Submissions', and 'Logout'. The main section is titled 'Upload Presentation' and includes a 'Presentation File * Required' field with a 'Choose File' button and the filename '2_Andy Jassy AgentCore.mp3'. Below this is a 'Presentation Title * Required' field with the text 'test for #2'. A 'Description (optional)' field is at the bottom with the placeholder 'Describe your presentation'.

The screenshot shows the 'My Submissions' page. It lists three submissions: 'test for #2' (Pending), 'Test for #1' (Processing...), and 'test of Pres #4' (Completed). Each submission entry includes the filename, upload date, and a 'Delete' button. The 'test for #2' entry also has a 'Pending' status badge, 'test for #1' has a 'Processing...' badge, and 'test of Pres #4' has a 'Completed' badge.

User-facing flow



The deliberate split: deterministic processing does the mechanical work, while the probabilistic agentic layer handles judgment. That separation is what makes the design easier to govern.


Let's Deep Dive into the App UI Frontend & Backend, processing, and Agentic layers and then review the AI-DLC behind it.

The Web App

Responsive Web App




Web



Mobile

Cognito Based User IAM



Groups:

- Users
- Administrators

Other Cognito Functions implemented

- User self signup
- Password reset
- Password recovery
- MFA not implemented in MVP

Audio File Upload

 **Presentation Coach**

[Administrator](#) [Upload](#) [Submissions](#) [Logout](#)

Upload Presentation

Presentation File * Required

Choose File

1_Introducing AgentCore.mp3

1_Introducing AgentCore.mp3 (2.45 MB)

Presentation Title * Required

test run

9 / 200

Description (optional)

Describe your presentation

0 / 2000

Upload

 © 2026 Geislersoft Presentation Coaching Platform

Features

- File picker
- Title
- Description
- Video File upload on roadmap

Processing

 Presentation Coach		Administrator	Upload	Submissions	Logout
My Submissions					
test 3		Pending			
File: 3_Agentcore-Introduction.mp3 AgentCore Introduct from AWS Uploaded: Jun 20, 2026, 02:37 PM		Delete			
test of another presentation		Completed			
File: 2_Andy lassy AgentCore.mp3 Uploaded: Jun 20, 2026, 11:12 AM		Delete			
Download Report					
test repeat		Completed			
File: 1_Introducing AgentCore.mp3 Uploaded: Jun 20, 2026, 09:07 AM		Delete			
Download Report					
test		Failed			
File: 1_Introducing AgentCore.mp3 Uploaded: Jun 19, 2026, 02:52 PM		Delete			

Features

- Multiple steps
- Realtime User status updates

Statuses

Start of Presentation Coach Process

Report test File: 1_Introducing AgentCore.mp3 Verifying the report output Uploaded: Jun 22, 2026, 02:05 PM		Pending	Delete
Report test File: 1_Introducing AgentCore.mp3 Verifying the report output Uploaded: Jun 22, 2026, 02:05 PM	Processing...	Processing	Delete
Report test File: 1_Introducing AgentCore.mp3 Verifying the report output Uploaded: Jun 22, 2026, 02:05 PM		Waiting	Delete
Report test File: 1_Introducing AgentCore.mp3 Verifying the report output Uploaded: Jun 22, 2026, 02:05 PM	Processing...	Evaluating	Delete

Pending: User uploaded unprocessed audio file

Processing: Transcription and prep

Waiting: Queued for Agentic Evaluation

Evaluation: Agentic analysis in progress

Statuses

new test File: 1_Introducing AgentCore.mp3 test of Exec summary Uploaded: Jun 22, 2026, 09:42 PM	Processing...	Generating Report	Delete
Report test File: 1_Introducing AgentCore.mp3 Verifying the report output Uploaded: Jun 22, 2026, 02:05 PM		Completed	Delete

End of “Happy Path” Presentation Coach Process

test File: 1_Introducing AgentCore.mp3 Uploaded: Jun 19, 2026, 02:52 PM		Failed	Delete
--	--	--------	--------

Generating Report:
Final report is being created

Completed: Report is ready for download

Failed: An error has halted processing

Report is Ready to Read

 Presentation Coach		Administrator	Upload	Submissions	Logout
My Submissions					
test 3 File: 3_AgentCore-Introduction.mp3 AgentCore Introduce from AWS Uploaded: Jun 20, 2026, 02:37 PM		Completed		Download Report	
test of another presentation File: 2_Andy.Jassy.AgentCore.mp3 Uploaded: Jun 20, 2026, 11:12 AM		Completed		Delete	
test repeat File: 1_Introducing AgentCore.mp3 Uploaded: Jun 20, 2026, 09:07 AM		Completed		Download Report	
test File: 1_Introducing AgentCore.mp3 Uploaded: Jun 19, 2026, 02:52 PM		Failed		Delete	

Features

- Current status of analysis
- Data entered by User displayed
- Download report link
- Delete link

The Presentation Coaching Report

Presentation Coaching Report

User: Michael Geiser (mgeiser@mgeiser.net)
Presentation Title: new test
Description: test of Exec summary
Presentation File: 1_Introducing AgentCore.mp3
Uploaded: June 22, 2026 21:42 ET

Executive Summary

This presentation is technically knowledgeable and genuinely engaging in tone, but it reads more like an enthusiastic product walkthrough than a confident, authoritative introduction — and that gap is what the 5.9 overall average reflects. That score signals a presenter with real substance who has not yet learned to carry that substance with full conviction. The content is credible, the energy is real, but the delivery patterns, structural looseness, and absence of persuasive architecture consistently blunt the impact of what is actually strong material.

The clearest strengths here are the technical depth and the instinct to make complex infrastructure concepts feel human. Delivery scores the highest of any dimension for good reason: the conversational warmth, inclusive second-person framing, and natural rhetorical questions in the memory section all create a sense of a knowledgeable guide rather than a lecturer reading slides. Technical Communication also scores well because the problem-first framing is genuinely effective — introducing each capability by naming the pain point it solves is sound instructional architecture. The honest acknowledgment of platform limitations and the proactive handling of the vendor-lock objection demonstrate a presenter who understands the audience's skepticism, and that kind of intellectual transparency builds real credibility with technical listeners.

The two highest-leverage improvements cut across nearly every lower-scoring dimension simultaneously. First, Executive Presence and Persuasion both suffer from the same root cause: hedging and casual filler language — "kind of," "sort of," "super super easy," "mishmash," "if you ask me" — that systematically undercuts the authority this presenter has clearly earned. Replacing that language with precise, confident phrasing would lift scores across Delivery, Executive Presence, and Persuasion in a single pass. Second, Audience Engagement and Structure both suffer from the absence of a narrative spine. The presentation is currently a feature inventory with a loose framing device; transforming it into a problem-solution arc with a concrete protagonist — even a hypothetical developer struggling to get an agent into production — would give the audience something to follow emotionally, not just intellectually. Structure also needs a clean verbal agenda at the opening and a proper 30-to-45-second closing recap, rather than the current 12-second pivot to the demo. For the next recording, focus on three things: open with one bold, evidence-anchored statement of impact instead of "Hey everyone," replace every instance of "kind of" and "sort of" with either a precise word or a clean pause, and close with a 40-second summary that names the three capabilities the audience should remember and tells them exactly what they are about to do in the hands-on session and why it matters.

Per-Dimension Detailed Feedback

Delivery (Score: 6.8/10.0)

Findings:

- [MEDIUM] Filler Words: Two instances of 'uh' were detected in the transcript (at approximately 159s and 241s). Additionally, conversational hedges like 'you know' (at ~38s), 'kind of' (used multiple times), and 'sort of' appear several times throughout the presentation, reducing overall polish.
→ Suggestion: Practice pausing silently instead of using filler words. Record yourself and flag each 'uh,' 'you know,' and 'kind of' — replace them with a deliberate 1-2 second pause to maintain authority and clarity.
- [MEDIUM] Pace: The presentation spans approximately 307 seconds (just over 5 minutes) and covers a dense amount of technical content. Analysis of word timing shows a largely consistent and moderately fast pace with few natural breaks between major topic transitions. Some segments, like the capabilities overview (179s–267s), rush through multiple distinct features without sufficient pacing variation.
→ Suggestion: Intentionally slow down when introducing a new concept or feature (e.g., identity, memory, gateways). Aim for a 10–20% pace reduction on key points to allow the audience to absorb information before moving on.
- [MEDIUM] Pauses: There are some natural pauses at sentence boundaries (e.g., after 'It's really cool.' at ~9.7s, and after 'Big one is memory.' at ~221s), but strategic mid-section pauses for emphasis are largely absent. Transitions between major capability sections (identity → authentication → tools → memory → gateways) lack deliberate pause moments that signal a shift in topic.
→ Suggestion: Insert 2–3 second intentional pauses before introducing each major capability (e.g., before 'Big one is memory'). These signpost pauses help the audience mentally prepare to receive new information and add dramatic effect to key points.
- [MEDIUM] Vocal Variety: The delivery is conversational and relatively smooth, but the transcript shows a somewhat uniform rhythm throughout. High-value phrases such as 'production-level solutions,' 'real business value,' and 'Big one is memory' could benefit from stronger tonal emphasis. The use of informal language ('mishmash of stuff,' 'fuss around,' 'nitty gritty') adds personality but may dilute authority when not paired with vocal emphasis.
→ Suggestion: Identify 4–5 key phrases per section and rehearse delivering them with elevated pitch or increased volume. For example, slow down and drop pitch on 'that's not gonna cut it at scale in production' to underscore importance.
- [LOW] Energy: The presenter demonstrates genuine enthusiasm, particularly evident in phrases like 'It's really cool,' 'that's pretty cool if you ask me,' and 'super, super easy to deal with.' The casual, conversational tone conveys authenticity and approachability throughout the 307-second presentation.
→ Suggestion: Channel current enthusiasm more strategically — save the highest energy moments for the most critical benefits (e.g., when explaining memory management or the framework-agnostic nature of


Features

Detailed feedback across 7 dimensions:

- Delivery
- Structure
- Executive presence
- Technical communication
- Audience engagement
- Pacing
- Persuasion

Draft: Report format will be worked on during 60 Day Path to Production.

CI/CD Pipelines




The screenshot shows the AWS CodePipeline console. On the left is a sidebar with navigation links: Developer Tools, CodePipeline, and Pipelines. The main area displays a list of pipelines. Each pipeline entry includes a radio button for selection, the pipeline name, its latest execution status (all are 'Succeeded'), the latest source revisions (pulling from GitHub), the latest execution started time, and a link to view the most recent executions.

	Name	Latest execution status	Latest source revisions	Latest execution started	Most recent executions
☐	prescoach-dev-kiro-eval-workflow-full-deploy	✔ Succeeded	GitHub – e187925d : fix: corecting erros on eval work unit tests	2 days ago	✔ ✔ ✔ ✔ ✔ View details
☐	prescoach-dev-kiro-eval-workflow-deploy	✔ Succeeded	GitHub – 5743dcde : exec summary changes	2 days ago	✔ ✔ ✔ ✔ ✔ View details
☐	prescoach-dev-kiro-eval-workflow-test	✔ Succeeded	GitHub – 358e4f6b : fix: use cross-region inference profile us.anthropic.claude-sonnet-4-6 (required...	7 days ago	✔ ✔ View details
☐	prescoach-dev-kiro-prep-workflow-full-deploy	✔ Succeeded	GitHub – 30c80d86 : fix: use existing uploads bucket for vector storage instead of nonexistent presc...	9 days ago	✔ ✔ ✔ ✔ ✔ View details
☐	prescoach-dev-kiro-prep-workflow-test	✔ Succeeded	GitHub – 358e4f6b : fix: use cross-region inference profile us.anthropic.claude-sonnet-4-6 (required...	7 days ago	✔ ✔ View details
☐	prescoach-dev-kiro-prep-workflow-deploy	✔ Succeeded	GitHub – e0380c9a : fix: status and processing layout	4 days ago	✔ ✔ ✔ ✔ ✔ View details
☐	prescoach-dev-kiro-frontend	✔ Succeeded	GitHub – 7840c001 : html: UI tweaks	4 days ago	✔ ✔ ✔ ✔ ✔ View details
☐	prescoach-dev-kiro-backend	✔ Succeeded	GitHub – e0380c9a : fix: status and processing layout	4 days ago	✔ ✔ ✔ ✔ View details
☐	prescoach-dev-kiro-full-deploy	✔ Succeeded	GitHub – f53e6401 : feat: add report download link — generate presigned S3 URL for completed submiss...	7 days ago	✔ ✔ ✔ View details

CICD Pipelines pull from Github and allow deploys at multiple logical unit levels:


- Frontend/Backend
- Audio/Video Processing
- Agentic AI Evaluation




Each Pipeline can deploy

- Tests only ("test")
- Code only ("deploy")
- Tests & code ("full-deploy")

Administration – Environment Variables



“Adminitration” menu only available to users with Administrator role in Amazon Cognito.

Two side-by-side screenshots of the 'Environment Variables' configuration modal. The left screenshot shows the form with fields for model IDs, idle timeout, and max concurrent evaluations. The right screenshot shows the dropdown menu for 'SESSION_SUPERVISOR_MODEL_ID' open, displaying a list of available models.

Environment Variables

SESSION_SUPERVISOR_MODEL_ID
Foundation model used by the Session Supervisor agent

Claude Sonnet 4.6 (US CRI)

COACHING_SUPERVISOR_MODEL_ID
Foundation model used by the Coaching Supervisor agent

Claude Sonnet 4.6 (US CRI)

EVALUATION_MODEL_ID
Foundation model used by the individual evaluation agents

Claude Sonnet 4.6 (US CRI)

IDLE_TIMEOUT_MINUTES
Minutes of inactivity before the ECS evaluation task exits

30

MAX_CONCURRENT_EVALUATIONS
Maximum number of submissions processed simultaneously

5

Cancel Save Changes

Environment Variables

SESSION_SUPERVISOR_MODEL_ID
Foundation model used by the Session Supervisor agent

Claude Sonnet 4.6 (US CRI)

Amazon Nova Premier (US CRI)

Amazon Nova Pro (US CRI)

Amazon Nova Pro (Single Region)

Amazon Nova Lite (US CRI)

Amazon Nova Lite (Single Region)

Amazon Nova Micro (US CRI)

Amazon Nova Micro (Single Region)

Claude Sonnet 4.6 (US CRI)

Claude Sonnet 4.6 (Single Region)

Claude Sonnet 4.5 (US CRI)

Claude Sonnet 4.5 (Single Region)

Claude Sonnet 4 (US CRI)

Claude Sonnet 4 (Single Region)

Claude Haiku 4.5 (US CRI)

Claude Haiku 4.5 (Single Region)

Claude Opus 4.6 (US CRI)

Claude Opus 4.6 (Single Region)

Claude Opus 4.7 (US CRI)

Claude Opus 4.7 (Single Region)

Claude Opus 4.8 (US CRI)

Features


Full control over all environment variables

- Models used
- Idle timeout for ECS tasks
- AI Concurrency limits
- IAM settings

App designed as single tenant but multi-tenant was considered in design

Tested and approved models only are allowed to be selected for the app.

Administration – Feature Flags




Features

Full control over Feature Flagged scope:

- **Video Processing:** User could upload video files to be processed in addition to audio files
- **Batch Processing:** Batch processing allows for lower cost asynchronous processing of presentations for vector stores.
- **Embeddings:** turns on and off vectorization of presentations.
- **Local Mode:** using a simplified (less resilient) agentic processing implementation

Observability and Monitoring – (eta: Aug 15, 2026)

Development Costs



Features


Total costs of the development

- AWS Account: \$6.59
- Kiro AI-DLC: 1355 Credits (\$27.10)
- LLM usage is not included in these costs

Report issue  Kiro Pro+ 1354.89 / 2000 updated just now

Github Repo Activity


Initial Build




<https://github.com/michaelgeiser/kiro-agentcore-app>

The Architecture


Web App SPA and Backend Architecture




Audio File Preparation Architecture




Evaluation with Agentic AI and Report Creation



Option: Replacing EventBridge Pipes with Lambdas can be explored




Multi-Region Ready if the Business Case Exits



DELIVERY MODEL

Build approach: AI-DLC turned the idea into deployable work units


The useful lesson is not that code was generated quickly. The useful lesson is that the work was decomposed into independently testable “bolts” with clear contracts.



Retrospective describes four work units, bolt-based interfaces, CI/CD pipelines, and zero-idle service selection.

RUNTIME FLOW

End-to-end processing is event-driven and observable




The handoff contract matters more than the tool choice. The upload service does not know how evaluation works. The evaluator does not care how the file arrived. That keeps each boundary testable and replaceable.

State transitions
Pending → **Processing** → **Completed**
or **Failed**

Failure path
retries → **HandleFailure** → **SNS + DLQ**

Architecture, viewed like a CTO would review it



Why this is a good MVP architecture

- Files never pass through Lambda
- Async processing gives back-pressure
- Feature flags allow runtime tuning
- Least-privilege IAM can be reviewed per work unit

What I would challenge before production

- Model quotas and throttling behavior
- Regional failure stance and user-pool strategy
- PII retention and report lifecycle policy
- Evaluation result explainability and traceability


Architecture summarizes services described in project overview, flow, and retrospective.

The agentic layer is composable by design

This is the part that justifies an agentic approach. The value comes from multiple independent lenses reviewing the same asset, not from one big prompt pretending to be a committee.

Agent definitions live in a JSON manifest and can be enabled or disabled without code changes.

The report generator assembles structured results into a 10-18 page PDF coaching report.



COST AND SCALE

The economics are exactly why serverless fits

For spiky, early-stage usage, the cost goal is not “cheapest at theoretical peak.” It is “do not pay for infrastructure while product-market fit is still loading” but will scale economically as usage takes off.

LLM selection must be evaluated and optimized. Costs and output quality differ for LLMs such as Claude Sonnet 4.6 vs Amazon Nova Pro and need to be evaluated.

Demo / PoC

10/month

 **\$0.35****Regular use**

500/month

 **\$18****Conference spike**

10,000/week

 **\$350****Quiet after spike**

following month

 **\$3.55**

Infrastructure only: LLM costs must be determined; cost and output will vary by LLM.

**Bedrock dominates
the bill at scale**

**Zero idle
for Lambda, API Gateway,
Step Functions, SQS, SNS,
Bedrock**

**S3 vector store
keeps MVP cost near zero**

**S3 vector store
keeps MVP cost near zero**

COST DECISION

S3 as the first vector store is the right kind of cheap

This is one of those choices where CTO discipline matters. For the MVP, S3 is not the “best vector database.” It is the best way to avoid paying for one before the workload proves it deserves one; it is **fit to purpose** engineering.

S3 JSON embeddings

Idle cost

\$0 storage only

Read pattern

Per-submission retrieval

Latency

~50-100 ms per file

Move when

MVP and variable load

OpenSearch Serverless / pgvector

~\$175/month minimum

Similarity search across corpora

~10 ms query path


>50K presentations/month or cross-submission search

Design win: vector-store-type and vector-store-endpoint are SSM parameters, so the upgrade path is operationally clean.

RESILIENCE

Operational resilience is solid for MVP and launch, but the regional stance is a decision

Current state is single-region in us-east-1 and multi-AZ compatible by design and construction. Single Region is acceptable for an MVP, but production requires an explicit answer to how much regional failure the business will tolerate. Multi-region is cost effective and does not require re-architecture work.



Low-friction upgrades to multi-region resilience: DynamoDB Global Tables, S3 CRR, second regional API Gateway, replicated Step Functions and Lambda stacks, CloudFront origin routing.

RISK REGISTER

The production-hardening backlog is manageable, not mysterious

The system is deployed and the end-to-end pipeline runs. The remaining work is application-level correctness plus the usual enterprise-grade hardening. Good news: no smoke machine required.

Known app fixes

Use Web Sockets instead of polling for Submissions page while a presentation is processing.

Security

Add additional guardrails, add admin approvals for new users, enhance WAF rules, define data retention policies.

Reliability

Refine throttling policy by user and entire system, ECS Spot interruption handling, DLQ message runbooks, retry budgets, test idempotency fully.

Observability

Add Correlation Id and additional logging to trace a submission across Cognito, API, S3, Step Functions, queues, ECS, Bedrock, and report generation.

FinOps

Refine budgets, Bedrock spend alarms, lifecycle policies, model-choice controls, chargeback by user or tenant.


Evaluate items and prioritize:
In this case the Security and Reliability issues have the highest priority

Known application issues come from project overview and source deck. Hardening items are CTO recommendations derived from the architecture.

WHERE I WOULD SPEND TIME

CTO decision radar

The architecture is good enough to advance. The discussion should move from “can this work?” to “what operating posture do we want?”



This slide distills CTO decisions from the supplied architecture, flow, cost, and current-state documents.

NEXT MOVES

A Practical 15/30/60 Day Path to Production

Do not boil the architecture ocean. Fix the known seams, prove the evaluator quality, then harden only the parts that have a user or revenue reason to exist.

15 days

Stabilize MVP

Observability and monitoring enhancement

Administrator Menus

Refine throttling policy

ECS Spot interruption handling,

Budget alarms on Bedrock usage

30 days

Prove quality

Golden set for coaching quality

Prompt and model versioning

Agent output schema enforcement

Report scoring calibration

60 days

Prepare production posture

Retention and PII controls

Other Responsible AI controls

RTO/RPO decisions

Quotas and load test

Tenant or user-level cost attribution

Exit criteria

A CTO should be able to trace one submission from upload to report, explain one agent's score, and predict the cost of the next 1,000 runs.

CLOSE

The CTO-to-CTO takeaway

PresCoach is a strong pattern because it combines three things that usually get separated: a useful product experience, an AWS-native cost model, and an AI-DLC delivery method that keeps architecture boundaries visible.

What is working

Clear work-unit boundaries, async flow, managed services, zero idle cost, composable agents.

What to fix

Retrieval, parser handling, traceability, hardening of AI and data controls.

What to decide

Regional posture, vector store timing, quality gates, model spend controls, production-readiness threshold.


Bottom line: advance it but make the next phase *boring* in the best possible enterprise way. Boring is where systems go to become dependable.

Appendix Material

Architecture Flow Diagrams


Upload and Storage Service — Architecture Diagram

prescoach-dev-kiro




Preparation Workflow — Architecture Diagram

prescoach-dev-kiro | Step Functions Standard Workflow | 10 Lambda Functions



Agentic Evaluation — LOCAL MODE Architecture

ECS Fargate Spot | Strands SDK Local Execution | Direct Bedrock InvokeModel



LOCAL MODE — Key Characteristics

- All agents run in-process inside a single ECS Fargate Spot container (no AgentCore deployment required)
- Strands SDK creates Agent objects locally — each evaluation tool calls bedrock:InvokeModel directly
- Session state is in-memory only (lost when task exits) — no persistent memory across task launches
- Cost-optimized: Fargate Spot (70% discount), 30-min idle timeout, concurrent processing of 5 submissions
- Cold start: ~45-90s (container pull + startup). Kept warm for 30 min after last message.

Cost Estimate

Per evaluation task (5 min): ~\$0.0007


Monthly @ 5/day: ~\$0.10


Monthly @ 100/day: ~\$2.10


(Bedrock model costs are additional)

Agentic Evaluation — AGENTCORE MODE Architecture

Amazon Bedrock AgentCore Runtime | Managed Agents | Session Memory | Auto-Scaling







Replication
S3 CRR
&
Dynamo Global
Tables

